



# HuntKTm: Hybrid Scheduling and Automatic Management for Efficient Kernel Execution on Modern GPUs\*

WENXUAN PAN, School of Computer Science and Engineering, Sun Yat-Sen University, Guangzhou, China

ZEJIA LIN, School of Computer Science and Engineering, Sun Yat-Sen University, Guangzhou, China

JIANGSU DU, School of Computer Science and Engineering, Sun Yat-Sen University, Guangzhou, China

XIANWEI ZHANG, School of Computer Science and Engineering, Sun Yat-Sen University, Guangzhou, China

Nowadays, Graphics Processing Units (GPUs) dominate in a wide spectrum of computing realms with massive parallel processing capabilities. However, as resources are continuously integrated into GPUs, traditional serial execution often leads to underutilization. Prior studies have shown that allowing multiple kernels to run concurrently and share GPU resources can effectively improve both resource utilization and system throughput. Nonetheless, existing methods for solely automating concurrent kernel execution either schedule kernels within individual applications (i.e., kernel-level) or enable task concurrency across multiple applications (i.e., task-level), thus leaving substantial GPU capacity underexploited. Moreover, they inevitably introduce new programming frameworks, which incur cumbersome manual efforts and further impose substantial programming burdens on developers.

To address these limitations, we propose HUNTKTm, a hybrid scheduling and automatic management method that cooperates kernel-level and task-level concurrency to enhance system throughput with minimal code modification. Specifically, HUNTKTm comprises a stream scheduler to assign kernels, a task scheduler to dispatch tasks onto GPUs, and a memory manager to reduce memory footprint. The stream scheduler applies a strategy to dispatch kernels to hardware queues and adopts a novel algorithm to remove redundant synchronizations in computational flow. Then, the task scheduler automatically issues tasks based on resource requirements and availability to support GPU sharing among uncooperative applications. Finally, the memory manager reduces memory footprint for tasks by limiting the lifetimes of memory objects, which thereby enables more tasks to execute simultaneously. Experimental results demonstrate that HUNTKTm improves system throughput by 33.2% over the existing state-of-the-art CASE framework on a single machine equipped with four NVIDIA A100 GPUs and reaches 13.8% higher application-level performance over Taskflow, even with lessened programming efforts.

**\*Extension of Conference Paper:** Zejia Lin, Zewei Mo, Xuanteng Huang, Xianwei Zhang, Yutong Lu. 2023. KeSCo: Compiler-based Kernel Scheduling for Multi-task GPU Applications. In *IEEE 41st International Conference on Computer Design, Washington DC, USA*. This extended version makes the following new contributions to the conference paper: 1. We observe that solely applying kernel scheduling, e.g., KeSCo, or task scheduling fails to fully leverage the potential of concurrent execution to enhance GPU resource utilization. 2. Based on this insight, we propose a hybrid scheduling strategy comprising a stream scheduler and a task scheduler. The strategy automatically schedules a task's kernels to different streams and dynamically dispatches tasks to suitable devices based on resource demands and availability. 3. To address the memory capacity bottleneck in hybrid scheduling, we introduce a memory manager that reduces the memory footprint of tasks using liveness analysis. 4. We conduct a more comprehensive evaluation to analyze the performance benefits of the hybrid scheduling strategy and memory management in both multi-task and single-task scenarios, offering insights into the underlying causes of the improvements. Experimental results demonstrate that HUNTKTm significantly improves system throughput and accelerates task execution over prior arts. This research is supported by the National Key R&D Program of China (Grant No. 2023YFB3002202), the NSFC grants (62472462, 62461146204). Xianwei Zhang and Jiangsu Du are the corresponding authors.

---

Authors' Contact Information: Wenxuan Pan, School of Computer Science and Engineering, Sun Yat-Sen University, Guangzhou, Guangdong, China; e-mail: panwx5@mail2.sysu.edu.cn; Zejia Lin, School of Computer Science and Engineering, Sun Yat-Sen University, Guangzhou, Guangdong, China; e-mail: linzj39@mail2.sysu.edu.cn; Jiangsu Du, School of Computer Science and Engineering, Sun Yat-Sen University, Guangzhou, Guangdong, China; e-mail: dujiangsu@mail.sysu.edu.cn; Xianwei Zhang, School of Computer Science and Engineering, Sun Yat-Sen University, Guangzhou, Guangdong, China; e-mail: zhangxw79@mail.sysu.edu.cn.



This work is licensed under a Creative Commons Attribution 4.0 International License.

© 2025 Copyright held by the owner/author(s).

ACM 1544-3973/2025/11-ART

<https://doi.org/10.1145/3774652>

CCS Concepts: • **Software and its engineering** → **Compilers; Runtime environments; Scheduling; Software performance; Massively parallel systems.**

Additional Key Words and Phrases: Concurrent kernel execution, Task scheduling, GPU memory management, Compiler, Runtime system

## 1 Introduction

In the last decade, Graphics Processing Units (GPUs) have been widely applied in a myriad of domains, owing to their excessive computation capability and high memory throughput. Advanced GPUs incorporate more resources than what a typical monolithic GPU task<sup>1</sup> necessitates and are thus frequently being underutilized, especially when executing single-kernel programs, which launch only one kernel at a time. To alleviate the underutilization issue, a plethora of approaches have been proposed, with representative schemes of concurrently executing sliced kernels [36, 53, 56, 57] and resource virtualization [12, 40, 43, 52].

As GPU applications become more complex, multi-kernel programs, originally consisting of concurrently executable kernels, are emerging across diverse realms. Compared to single-kernel programs, multi-kernel ones can leverage various GPU streams and synchronization events to parallelize kernel executions to efficiently shorten execution time. Such an optimization requires developers to correctly analyze dependencies between kernels and then rearrange kernels in task queues to strike load balance and minimize synchronization cost. Without a doubt, considerable programming efforts should be made to obtain bug-free and highly performant codes, particularly for increasingly complicated programs and architectures. To address the issue, a bunch of designs have been recently presented to automate inter-kernel concurrency of GPU applications. A GrCUDA-based [29] runtime approach GrSched [38] applies a virtual machine, exempting developers from the need to explicitly claim kernel dependencies. But compared to expertise-based optimizations, GrSched introduces serious performance downgrade due to the overheads of runtime scheduling. Instead, Taskflow [15] proposes a new heterogeneous programming framework to automate inter-kernel concurrency optimization. It harnesses cudaGraph [33] to reduce the overheads of fragmented kernel launches. Nevertheless, such a method requires developers to grasp a new programming model and manually specify kernel dependencies, inevitably raising coding difficulty. Although executing kernels concurrently within an application greatly enhances GPU resource utilization, the improvement is constrained by the limited number of concurrent kernels.

Another direction to address resource underutilization is sharing GPU among tasks. In concurrent task execution, time-sharing is most commonly used, which allocates time slices for kernels to perform computations in turn. Various researches have been performed to optimize task scheduling with a time-sharing model [6, 13, 14]. These methods aim to ensure quality of service (QoS) in concurrent execution but have limited benefits on utilization and overall throughput. NVIDIA Multi-Process Service (MPS) [28] enables kernels from different processes to execute on the same device, implementing space-sharing by partitioning resources according to user-defined computing resource thresholds. Building upon MPS, GSlice [10] introduces a self-learning resource allocation algorithm to assign suitable resources for DNN inference tasks. SchedGPU [44], a runtime task scheduling system, seeks to maximize task concurrency on a single GPU while avoiding out-of-memory (OOM) issues. To facilitate the concurrent execution of more tasks, researchers have expanded the scheduling scope to encompass multiple GPUs. In multi-GPU systems, task scheduling requires explicitly specifying the target device based on the resource requirements of tasks, imposing additional programming burdens on users. CASE [4] reduces such manual efforts by statically analyzing task resource requirements, and then dynamically assigning tasks to appropriate devices according to available GPU resources. However, memory capacity emerges as a performance bottleneck that limits the ability to launch additional tasks. Furthermore, as the number of tasks

---

<sup>1</sup>Here a task refers to an independently executing program that consists of one or more kernels.

increases, the interference between tasks becomes significant and cannot be overlooked. Therefore, relying solely on kernel scheduling or task scheduling is insufficient to fully utilize the available GPU resources.

To maximize resource utilization in GPU systems, we propose HUNTKTm, a hybrid scheduling and automatic management method that cooperates kernel-level and task-level concurrency to facilitate efficient GPU execution. HUNTKTm is comprised of a *stream scheduler* and a *task scheduler* to combine concurrent kernel and task scheduling, and a *memory manager* for memory footprint reduction. Specifically, *stream scheduler* automatically identifies data dependencies and places kernels into different streams concerning load balance and synchronization cost. *Task scheduler* analyzes resource requirements of tasks and available system resources, and subsequently performs efficient, memory-safe device dispatch across GPUs. *Memory manager* conducts liveness analysis on multi-kernel programs and facilitates memory reuse by scheduling memory allocation and deallocation instructions. The source code of HUNTKTm is available at <https://github.com/Gemini321/HuntKTm>.

In summary, the contributions of this paper are as follows:

- We highlight the inadequate performance enhancement and programming weakness of prior arts in solely improving intra-application or inter-application concurrent execution for multi-kernel programs.
- With the insight, we propose hybrid scheduling by encompassing a stream scheduler to exploit kernel-level concurrency for multi-kernel programs, and a task scheduler to analyze resource requirements and dispatch tasks to appropriate devices automatically.
- To widen scheduling space, we further present a memory manager based on memory liveness analysis to reduce task memory footprint. The manager alleviates the memory capacity bottleneck and accounts for launching sufficient tasks to saturate system resources.
- Experimental results show that HUNTKTm effectively improves execution performance and resource utilization of GPUs. Compared with the state-of-the-art scheduling framework CASE, HUNTKTm achieves an average of 33.2% throughput improvements. Additionally, HUNTKTm delivers 13.8% application-level performance improvement over Taskflow.

## 2 Background

This section mainly introduces GPU concurrency models and programming abstractions, and outlines memory management techniques focused on lifetime-aware memory allocation and reuse, which lays the foundation for subsequent system design and optimizations.

### 2.1 GPU Concurrent Execution

Designed for massively parallel computation, modern GPUs are typically equipped with many *streaming multi-processors* (SMs), each of which has hundreds of computing cores and can simultaneously execute up to thousands of threads. In many scenarios, one single kernel cannot fully utilize GPU hardware resources, thus causing a great waste of computation capabilities and low performance [54]. To alleviate such a problem, concurrent kernel execution (CKE) has been widely supported by vendors to parallelize inter-kernel execution on hardware components. It issues operations in multiple software queues (called *streams* in CUDA [33]), which are mapped onto different hardware queues and processed concurrently if the demanded resources, typically SMs, are sufficient. The multi-kernel workloads provide a perfect scenario for implementing CKE, as they have independent kernels that are ready to execute concurrently. However, the concurrency capability of CKE is limited by the number of independent kernels in a single program, making it challenging to saturate the available hardware resources.

Another feasible way to achieve higher GPU utilization is applying task-level concurrent execution among independent workloads. When multiple tasks run simultaneously, thread blocks from different tasks are able to occupy more SMs. Moreover, a GPU can swiftly switch to another task if one is stalled on memory access, effectively overlapping computation and communication. NVIDIA multi-process service (MPS) [28] provides




Fig. 1. Execution of concurrent kernels on a GPU [33].

a software mechanism that enables multiple processes to co-execute on the same device in a space-sharing fashion, which reduces context-switching overhead and improves GPU utilization across processes. Since Ampere architecture, NVIDIA multi-instance GPU (MIG) [30] splits a single GPU into multiple instances for clients, which provides isolated compute and memory resources for each instance. This technique enables predictable performance and improved resource utilization in multi-tenant environments.

## 2.2 CKE Programming in CUDA

To facilitate efficient computation with CKE, many popular GPU programming models [2, 11, 33] offer a series of concurrency APIs, here we take CUDA as an example. A data flow graph (DFG) needs to be constructed correctly first to help schedule the executions. The DFG is further divided into multiple levels such that kernels from the same level have no data dependence. Then developers need to create multiple *CUDA streams*, and issue kernels on different streams to co-execute on GPUs. To ensure the execution order of data-dependent kernels across streams, *CUDA events* are inserted after a kernel's predecessors to track their completion. These events are subsequently synchronized before the launch of the dependent kernel. Synchronization between kernels and asynchronous data copy follows a similar pattern.

Figure 1 shows an example of three concurrent tasks sharing a GPU. Kernels *B* and *C* are mutually independent, and they both depend on kernel *A*. After kernel *A* finishes, kernels *B* and *C* are issued on different streams and executed simultaneously on different SMs. At the same time, an asynchronous copy is proceeding on the copy engine, which is a complementary hardware resource for SMs. Therefore, the computation of two kernels and data transfer are overlapped, helping utilize the abundant resources of GPUs. However, developers need to properly scrutinize the complex dependencies, schedule kernels in streams, and generate synchronization barriers in CKE programming. Such code reorganization incurs tremendous manual effort and is also error-prone.

## 2.3 GPU Memory Management

Memory management is crucial for effective resource utilization in GPU computing [18, 21, 23, 49]. A memory object is a contiguous memory region owned by a variable, constituting a fundamental memory management unit. Basic APIs like `cudaMalloc` and `cudaFree` handle memory allocation and deallocation of memory objects, requiring device synchronization and lacking flexibility for variable-sized data. More advanced APIs like `cudaMallocAsync` and `cudaMallocManaged` are designed to enhance performance and usability. Beyond these APIs, memory reuse has been a primary focus in memory management optimization. This technique minimizes the footprint by reallocating unused memory, allowing more tasks to run simultaneously within limited memory capacity. A feasible way to reuse memory is by limiting the lifetimes of memory objects through allocation and deallocation scheduling. The lifetime of a memory object spans from its allocation to its deallocation. If the lifetimes of two memory objects do not overlap, the memory released by one can be reallocated to the other, enabling memory reuse between the two objects. A live range denotes the period between the initial and final access of a memory object, representing the shortest duration during which the object is utilized. By identifying

the live ranges of memory objects and tailoring the lifetimes to the live ranges, non-overlapping memory regions can be reused to reduce the overall memory consumption of tasks.

### 3 Related Work

#### 3.1 Concurrent Kernel Execution

CKE has been widely studied in fine granularity. Elastic kernel [36] slices kernels into multiple small ones and deploys them on different SM to speed up. Similarly, OpenMP [35] is leveraged to decompose kernels into multiple tasks in Junggler [3] and schedule tasks with dependencies via runtime mechanism. Also, Pagoda [59] concurrently executes narrow tasks at warp-level by virtualizing GPU resources and issues kernels when required resources are available. ROSGM [22] leverages stream priorities in the Robot Operating System to dynamically switch kernel scheduling strategies across different application scenarios. To reduce runtime profiling overhead, a Fisher feature selection-based method [47] is employed to classify kernels and collocate those with complementary characteristics. cCUDA [46] performs online profiling and ranking for kernels, and employs kernel slicing to maximize computation overlap. FlexSched [25] further enables dynamic resource allocation and preemptive scheduling during concurrent kernel execution using persistent kernels. Specifically, cCUDA and FlexSched assume that all kernels are ready before scheduling without data dependencies, focusing on selecting optimal subsets of concurrent kernels and devising strategies for resource allocation among them. Addressing concurrency among data-dependent kernels within an application, Taskflow [15] wraps GPU programming model APIs and implements a static scheduler in the framework. A GrCUDA [29]-based runtime scheduler [38] eases the prototyping of parallel applications. Distinguished from those prior arts, our proposed HUNTKTm aims to statically automate kernel scheduling in multi-kernel programs, achieving much better performance with reduced programming burden.

#### 3.2 Task Scheduling

A bunch of task schedulers have been proposed to optimize concurrent task execution. On the hardware level, new APIs are introduced in [41] to heterogeneous system architecture (HSA) for applications specifying task priority. Chimera [37] extends the SM scheduler to estimate the cost of kernel preemption to minimize the overhead. Similarly, the command buffer and status table are further embedded in the SM scheduler [50] to minimize the overhead for prioritized tasks. On the software level, FLEP [55] leverages a compiler-runtime system to control task preemption at the kernel level. EffiSha [5] schedules kernels at thread-block level dynamically with an online cost model. ElasticBatch[42], gputlet[9], Paris&ELSA[19] introduce innovative partitioning and scheduling algorithms designed for efficiently distributing inference requests across GPUs with MIG enabled. SchedGPU [44] co-locates applications on a device in a memory-safe manner through a dedicated runtime system. CASE [4] introduces a novel compiler-based approach for scheduling uncooperative tasks over a multi-GPU system, which shares some similarities to HUNTKTm regarding retrieving resource requirements in a lazy runtime. While prior arts consider resource requirements as static features, HUNTKTm integrates with a memory management strategy for alleviating the memory bottleneck in task co-execution scenarios.

#### 3.3 GPU Memory Optimization

The efficient utilization of limited GPU memory has been extensively studied across various scenarios. Techniques such as swapping [18, 23, 24, 58], recomputation [7, 17, 49], compression [16, 26, 45], and reusing [21, 39, 48] have been widely adopted for memory optimization. DeepUM [18] proposes a correlation prefetch technique to hide significant overhead due to unified memory page faults. DELTA [51] and ATP [8] combine both swapping and recomputation to achieve lower memory consumption and higher throughput in DNN training. SMC [26] selectively compresses read-only pages to enable memory oversubscription while avoiding severe decompression




Fig. 2. SM occupancy of co-running two memory-intensive applications with concurrent kernel execution (*multi-stream*), concurrent task execution (*multi-task*) and hybrid execution (*hybrid*).




Fig. 3. SM occupancy of running six applications, where *hybrid* launches two jobs in a batch and *hybrid w/ mem.* with reduced memory usage launches three jobs simultaneously.

cost. Targeting memory reusing, frameworks like PyTorch [39] and TensorFlow XLA [48] pre-allocate a memory pool before execution and adopt in-place operations to reuse memory spaces of input and output tensors. Occamy [21] analyzes tensor liveness among DNN and applies kernel fusion to eliminate redundant intermediate tensors. However, these reuse methods focus on highly structured deep learning applications and cannot be adapted for general GPU programs.

#### 4 Motivation

In this section, we provide intuitive examples to demonstrate the benefits of increasing kernel-level and task-level concurrency over GPU space-sharing scenarios. Then we discuss how to achieve this goal through automatic stream and task scheduling, and further memory management.

##### 4.1 Insufficiency of Only Kernel-level and Task-level Concurrency

As discussed in Section 2.1, space-sharing schemes are commonly applied to improve GPU hardware utilization. However, we observe that computing resources can remain underutilized even with GPU utilization sustained at 100% in some multi-task scenarios. Here we use SM occupancy, obtained through a lightweight GPU monitoring tool DCGM [32], as a metric to evaluate the utilization of GPU computing resources during kernel execution. Figure 2 illustrates the SM occupancy of running two memory-intensive applications M2, composed of several activation and reduction kernels from NVIDIA FasterTransformer [31], under three different concurrency schemes. *Multi-stream* issues multiple kernels without data dependencies from an application through several hardware queues simultaneously, while two applications execute serially. *Multi-task* co-execute two tasks in the same device with MPS being enabled. Even with 100% GPU utilization during execution, both *multi-stream* and *multi-task* achieve less than 10% SM occupancy. Such low occupancy indicates that relying solely on concurrent kernel execution or task execution can not fully exploit GPU resources. By combining *multi-stream* and *multi-task*, *hybrid* allows for the simultaneous execution of more kernels from different tasks, enhancing resource utilization and accelerating computation. Therefore, *hybrid* achieves an SM occupancy of up to 22%, offering a throughput improvement of 73.1% compared to *multi-stream* and 79.7% compared to *multi-task*. The result reveals significant potential for combining concurrent kernel and task execution in multi-task scenarios.

##### 4.2 Memory Capacity Bottleneck in Concurrent Task Execution

With combined concurrent kernel and task execution, the memory usage becomes a primary constraint on task-level concurrency, as GPU out-of-memory (OOM) can cause program crashes. Figure 3 shows the SM occupancy curves for six multi-stream applications running under MPS, both without memory management

Table 1. Summary of various concurrent schemes.

Scheme	D.A. <sup>a</sup>	C.M. <sup>b</sup>	M.M. <sup>c</sup>	T.S. <sup>d</sup>	N.P.F. <sup>e</sup>	Language
Serial	✗	✗	✗	✗	✓	C++
Taskflow [15]	✗	✓	✗	✗	✗	C++
GrSched [38]	✓	✓	✗	✗	✗	Python
PyTorch [39]	✓	✓	✓	✗	✗	Python
CASE [4]	✗	✗	✗	✓	✓	C++
HUNTKTm	✓	✓	✓	✓	✓	C++

**Notes:**

- <sup>a</sup> Automatic Dependency Analysis
- <sup>b</sup> Automatic Concurrency Management
- <sup>c</sup> Automatic Memory Management
- <sup>d</sup> Automatic Task Scheduling
- <sup>e</sup> No New Programming Framework

(*hybrid*) and with memory management (*hybrid w/ mem.*). The peak memory consumption is defined as the maximum instantaneous GPU memory during execution. In *hybrid w/ mem.*, we manually schedule the allocation and deallocation instructions in M2, decreasing M2’s peak memory consumption from 17.6 GB to 11.2 GB by reusing non-overlapping memory objects. Without memory management, only two applications could share a GPU with 40 GB memory, requiring three separate launches to complete the computation. With reduced memory footprint, three applications could be launched altogether. Running more tasks concurrently on a single device allows additional kernels from different tasks to saturate idle computational resources. Moreover, the result reveals that even with an increased number of concurrent tasks, *hybrid w/ mem.* exhibits minimal growth in data initialization and transfer time before kernel execution. This can be attributed to the overlap of communication time for certain tasks with the computation of others, along with the parallel execution of host operations across tasks. As a result, the enhanced task concurrency enabled by efficient memory management leads to an 18.2% improvement in system throughput. This indicates that memory capacity becomes increasingly critical for collocating more tasks on a single device and improving system throughput under hybrid scheduling.

### 4.3 High Programming Burden in GPU Management

With no doubt, considerable efforts and expert knowledge are needed to write CKE codes, which significantly raises the programming barrier and is error-prone. Table 1 summarizes the characteristics of various concurrency optimization schemes in terms of programming efforts. To reduce the programming complexity of CKE, approaches like Taskflow [15] and a GrCUDA-based [29] scheduler (aliased as GrSched for ease of reference) [38] have been proposed to craft new programming frameworks by extending CUDA’s API for stream management and synchronization. Taskflow demands explicitly specifying dependencies through its APIs, while GrSched introduces a DSL embedded in Python to support automatic analysis and scheduling. Both schedulers require thorough refactoring of source code, which becomes another programming burden for users. In particular, errors of manually specifying dependencies in Taskflow can lead to incorrect computation results that are often hard to detect and debug, further increasing the risk and effort in development. Similar to CKE, memory management is impractical for programmers, as it requires detailed knowledge of memory objects and precise control over allocation and deallocation. Empirically, misordered memory operations and missing synchronizations are common in manual memory management, often resulting in unpredictable and hard-to-reproduce errors. While frameworks like PyTorch [39] and Tensorflow [1] offer dynamic memory management support, they are designed for deep learning rather than general-purpose GPU computing and require additional efforts to program within the framework.

Problems become more sophisticated when concurrent execution extends to multi-GPU systems. Users must specify target devices for applications before running and carefully control execution timing and order to prevent crashes from OOM errors and performance degradation. CASE [4] introduces a runtime system that transparently schedules tasks to appropriate devices without requiring a new programming framework. However, CASE lacks in-depth program analysis and optimization, leaving the burden of kernel concurrency and memory management on programmers. Therefore, there remains the urgency for a method that enables efficient kernel execution and task scheduling in multi-GPU systems with lowered coding efforts.




Fig. 4. Overview of HUNTKTM, which consists of three components: *stream scheduler*, *task scheduler* and *memory manager*. *Stream scheduler* assigns kernels from multi-kernel programs to concurrent queues. *Task scheduler* evaluates the resource requirements of transformed programs and dispatches them to appropriate devices. *Memory manager* optimizes memory allocation and deallocation to reduce the memory footprint of tasks.

## 5 Design

This section presents our proposed design HUNTKTM, which incorporates hybrid scheduling and memory management for efficient task and kernel execution. We first give the overview and then elaborate on each module.

### 5.1 System Overview

HUNTKTM is a GPU kernel scheduling framework that supports efficient kernel execution through a combination of kernel-level and task-level concurrency scheduling along with automatic memory management. As shown in Figure 4, HUNTKTM consists of three key components: *stream scheduler*, *task scheduler*, and *memory manager*. *Stream scheduler* focuses on intra-task kernel scheduling by distributing kernels with data dependencies to different streams at compile time. It takes the source code of multi-kernel programs as input, analyzes inter-kernel data dependencies, and generates high-performance multi-stream programs to enable concurrent kernel execution within a task. Then, *task scheduler* combines compile-time and runtime information to evaluate the resource requirements of the program and dynamically collocates tasks on appropriate devices based on the available resources in multi-GPU systems. Cooperating *stream scheduler* and *task scheduler* automates the hybrid concurrent execution of kernels and tasks.

To address the memory bottleneck in hybrid scheduling, *memory manager* performs memory lifetime management during compilation. It processes the stream graph from *stream scheduler*, applying a novel analysis algorithm to identify the live range of each memory object, which corresponds to the actively used periods. By scheduling allocation and deallocation instructions, *memory manager* effectively shortens the lifetime of memory objects and reuses the objects with non-overlapping lifetimes, thereby reducing the peak memory usage. In summary, HUNTKTM maximizes concurrency through hybrid scheduling and memory optimization, ultimately improving GPU resource utilization and system throughput.

### 5.2 Stream Scheduler

To achieve kernel-level concurrency, the *stream scheduler* relies on lightweight code modifications to help construct the DFG (Data Flow Graph) of kernels, then schedules the kernels across multiple streams while ensuring correct execution order through inter-stream synchronization instructions. Figure 5 shows the overall structure of the *stream scheduler*, which consists of three main components: *DFG constructor*, *kernel distributor*,




Fig. 5. Workflow of the *stream scheduler*, which transforms serial source code into a stream graph with efficient kernel-level parallelism.




Fig. 6. An example of stream scheduling for input DFG. (a) A DFG organized into three levels to schedule ten kernels onto three available streams. (b) Scheduling result of *kernel distributor* and *synchronization generator* for DFG.

and *synchronization generator*. The input of *stream scheduler* is the source code of a task, containing a series of kernels coded in the form of sequential execution. The *DFG constructor* analyzes each kernel’s inputs and outputs to build a DFG based on data dependencies. The *kernel distributor* assigns kernels to multiple streams based on the analyzed DFG, and then the *synchronization generator* creates synchronization instructions between kernels in different streams to ensure the program executes correctly. The final output of the stream scheduler is a stream graph that contains information about the multi-stream execution, which can be further optimized by the following transform pass.

To construct the DFG accurately, *DFG constructor* first distinguishes read-only and writable kernel parameters by introducing a lightweight code modification: a constant is inserted before each kernel’s input parameters to indicate the number of following writable parameters. This modification enables the DFG constructor to identify potential write conflicts without analyzing the kernel’s source code. When building the DFG, we consider three types of data dependencies: Read-After-Write (RAW), Write-After-Read (WAR), and Write-After-Write (WAW), which are common in real-world HPC and AI applications. Directly constructing the kernels’ dependency graph from the complex control flow incurs significant overhead. To reduce the computational cost, we build the DFG by leveraging the sequential order of kernel launch. Kernels are iterated in reverse order and a breadth-first search algorithm is adopted to identify each kernel’s direct predecessors until all kernels are traversed. We determine inter-kernel dependencies based on whether different kernels access the same data object, with the condition that at least one of these accesses involves a write operation. For example, if kernel A precedes kernel B in the execution order, and both kernels mark a shared variable *data* as writable, then a WAW dependency is recorded from A to B. If *data* is read-only in both kernels, no dependency is added. Furthermore, pointer arguments derived from the same base address are treated to access the same data. This unification ensures that aliasing caused by pointer arithmetic does not result in missing dependencies.

When the DFG is constructed, *kernel distributor* assigns kernels to GPU streams to enable kernel-level concurrency. The process begins by leveling the DFG so that kernels in the same level have no mutual data dependencies. Then, the kernels in the DFG are assigned to different streams level by level. We define the *preferred predecessor set (PP-Set)* of an unscheduled kernel as the subset of its predecessors that are located at the end of streams. Kernels in the same level are first sorted by the size of their *PP-Set*, and those with smaller *PP-Set* are scheduled first to minimize cross-stream synchronization. When scheduling kernels to different streams, *kernel distributor* follows a set of rules: ① Kernels without any predecessor are evenly distributed across streams in a round-robin fashion. ② Kernels with a single predecessor are assigned to the same stream as that predecessor. ③ Kernels with multiple predecessors are assigned to the stream of the predecessor in their *PP-Set* that has the fewest unscheduled successors. This heuristic algorithm ensures that kernels are scheduled as early as possible after their predecessors while balancing stream workloads and reducing synchronization overhead.




Fig. 7. Structure of *task scheduler* in HUNTKM, which operates in two phases: offline and online. *Resource analyzer* retrieves the resource requirements from the input stream graph and incorporates the results into the program during compilation. In the online phase, each task launches a *lazy engine* that delays the execution of GPU operations and communicates with *task scheduler*. *Task dispatcher* maintains information about each GPU and is responsible for dispatching tasks to devices that satisfy the resource requirements.

Here we exemplify the above steps with regard to the given DFG in Figure 6(a) and the corresponding kernel distribution strategy in Figure 6(b). In level 1, kernels are evenly assigned to different streams according to Rule ❶. In level 2, kernel *F*, having the smallest *PP-Set*, is scheduled first and placed after kernel *C* in accordance with Rule ❷. After updating the *PP-Sets*, kernel *E* now has a smaller set than kernel *D*, as kernel *C* is no longer at the end of its stream, and is thus arranged after kernel *A* by Rule ❸. Finally, kernel *D* is placed after kernel *B* following Rule ❹. In level 3, we repeat the process and schedule them in the order of kernel *H*, *I*, and *J*, which are all placed after their *preferred predecessor*. Lastly, kernel *G* can choose from stream 1 and 3, where its predecessors are seated, and are randomly inserted in stream 3, as shown in Figure 6(a).

After scheduling kernels in asynchronous streams, *synchronization generator* comes into play to ensure the correctness of the execution order. A naive approach creates barriers whenever data dependence exists. However, some barriers are redundant and may cause performance overhead. To tackle this issue, a pruning algorithm is proposed based on the implicit synchronizations brought by the transitivity of dependency and serial execution of kernels in the same stream. When finished, the barriers are pruned to the minimum.

The *synchronization generator* traverses the kernels in each stream and works in three steps, suppose it is working on kernel *K*. In step ❶, it creates barriers for each of *K*'s predecessors that do not share the stream with *K*. In step ❷, it checks *K*'s predecessors in each stream, and reserves only the synchronization issued from the last predecessor in that stream. In step ❸, it enumerates kernels before *K* in the same stream, say *T*. If *K* and *T*'s predecessor share the same stream, and *K*'s predecessor is executed before *T*'s, *K* is then implicitly synchronized by *T* and *T*'s predecessor. Therefore, *K*'s barrier to that predecessor is safe to be removed. A full analysis of the complete DFG helps eliminate these redundant barriers correctly. In runtime analysis of GrSched, such elimination is infeasible due to the lack of a global view of the graph. The example of Figure 6(b) shows the barriers generated in solid lines and the removed barriers in dashed lines. *Synchronization generator* scans stream 1 and creates kernel *E* and *I*'s barriers by step ❶. The same is true for kernel *D* in stream 2 and kernel *J* in 3. For kernel *G*, step ❸ detects its implicit synchronization with kernel *A* by the execution order of *A* → *E* → *J*, so the barrier is removed.

### 5.3 Task Scheduler

*Task scheduler* comprises three components: *resource analyzer*, *lazy engine*, and *task dispatcher*. *Resource analyzer* operates during the compilation phase, analyzing each kernel's launch configuration and the size of memory objects. Then *resource analyzer* aggregates the computing and memory resource requirements for each task. *Lazy engine* collects the resource information that cannot be determined during static compilation at runtime. It defers

**Algorithm 1:** Resource-aware Task Scheduling Algorithm

---

**Input:** List of available GPUs  $GPUList$ , queue of pending tasks  $pendingQueue$ , task to be scheduled  $task$

**Output:** Scheduled target GPU  $targetGPU$

```

1 Function TaskSchedule( $GPUList$ ,  $pendingQueue$ ,  $task$ ):
2    $targetGPU \leftarrow None$ ,  $maxAvailSM \leftarrow 0$ ;
3   for  $GPU \in GPUList$  do
4     if  $task.memReq < GPU.freeMem$  and  $numStream < GPU.freeQueue$  then
5        $SMInThread \leftarrow (GPU.inUseThread + task.threadReq)/GPU.threadPerSM;$ 
6        $SMInReg \leftarrow (GPU.inUseReg + task.regReq)/GPU.regPerSM;$ 
7        $SMInSMem \leftarrow (GPU.inUseSMem + task.smemReq)/GPU.smemPerSM;$ 
8        $availSM \leftarrow GPU.numSM - max(SMInThread, SMInReg, SMInSMem);$ 
9       if  $availSM > maxAvailSM$  then
10         $targetGPU \leftarrow GPU;$ 
11         $maxAvailSM \leftarrow availSM;$ 
12      end
13    end
14  end
15  if  $targetGPU$  is not None then
16     $targetGPU.addGraphAndUpdateResource(task);$ 
17  else
18     $pendingQueue.push(task);$ 
19  end
20  return  $targetGPU;$ 
21 end

```

---

GPU-related operations when necessary, ensuring flexibility and adaptability to dynamic resource conditions. *Task dispatcher* integrates the task requirements provided by the *lazy engine* with the realtime system resource usage to select suitable devices for tasks. Together, these components enable *task scheduler* to dynamically and efficiently co-execute multiple tasks among GPUs.

To facilitate resource-aware task scheduling during compilation, *resource analyzer* extracts resource requirements, such as memory usage and number of threads, from the source code. In addition, the analyzer leverages vendor-provided compiler (e.g., nvcc [27]) to obtain the number of registers and the amount of shared memory required by each kernel. *Lazy engine* estimates the computing resources needed for each stream by using the resource requirements of the first kernel launched within the stream. The total computing requirement of the task is then represented by aggregating the requirements of all streams. For memory resource, *lazy engine* employs def-use analysis for memory objects to identify objects associated with each kernel and determines their sizes from the memory allocation instructions. Notably, some resource requirements depend on input size and cannot be determined at compile time. In such cases, these requirements are captured during runtime by intercepting CUDA calls through *lazy engine*. Meanwhile, *resource analyzer* inserts `cudaTaskSchedule` at points where resource requirements are fully determined, enabling the calculation of maximum memory and computing resources required for task scheduling. However, static analysis cannot derive certain information due to function encapsulation or complex control flow. If resource requirements remain undefined before the first kernel launch, `cudaTaskScheduleLazy` is inserted before each kernel launch to determine the execution device based on the currently requested resources instead of the total resource requirements.

After static resource analysis, *task scheduler* dynamically dispatches tasks to appropriate devices during runtime. Since the computing device remains undefined before task scheduling, *lazy engine* intercepts and delays GPU-related operations such as memory allocation, deallocation, and data transfers. When a program reaches the inserted scheduling instructions, *lazy engine* predicts the resource requirements based on the parameters of intercepted operations and the information provided by *resource analyzer*, then forwards the requirements to *task dispatcher*. And, *task dispatcher* schedules tasks based on resource requirements and returns the target device ID when scheduling finishes. Once scheduled to a specific device, *lazy engine* executes the intercepted operations in sequence and launches kernels only after all operations are completed. Compared to offline profiling or static analysis, lazy execution allows HUNT<sup>TM</sup> to determine each task’s resource requirements at runtime, providing accurate information for task scheduling without severe profiling overhead.

To mitigate the overhead caused by frequent memory allocation and deallocation during execution, *task scheduler* initializes a memory pool before the first allocation, with its size being determined by the predicted memory footprint. If the free memory in the pool is sufficient to satisfy the allocation request, the system directly returns the pre-allocated memory without invoking costly system calls. Upon deallocation, the runtime system retains the memory for reuse in subsequent allocation requests. The memory in the pool is fully released only when the application exits.

After collecting resource requirements, *task dispatcher* schedules tasks to appropriate devices based on resource information. The scheduling policy, detailed in Algorithm 1, takes as input the available GPUs, the pending task queue, and the task to be scheduled along with its resource requirements predicted by *lazy engine*. The algorithm iterates over available GPUs and selects one with sufficient free memory and adequate hardware queues (line 4 ~ 5). Since each SM within a GPU functions as an independent computational unit with various resources, the number of available SMs is used to represent the GPU’s computational capacity. *Task dispatcher* estimates available SMs from three perspectives: threads, registers, and shared memory, and prioritizes the GPU with the most available SMs (line 6 ~ 9). This approach prevents application failures due to memory shortages and alleviates resource conflicts by distinguishing between various resource utilization types. As a result, computational load is more uniformly balanced across multiple GPUs. If no GPU meets both the memory and hardware queue requirements, *task dispatcher* suspends the task request in a queue and retries scheduling whenever resources get released.

#### 5.4 Memory Manager

To address the memory capacity bottleneck during hybrid scheduling, we design a memory lifetime management method for multi-stream programs and integrate the method into *memory manager*. *Memory manager* takes the stream graph generated by *stream scheduler* as input. It first performs data flow analysis on each memory object to identify the live range of the objects. Based on the analysis results, *memory manager* schedules GPU memory allocation, deallocation, and other memory manipulation instructions to shorten memory objects’ lifetimes to their live ranges during compilation, thus reusing memory regions among objects with non-overlapping lifetimes. Finally, the reduced peak memory usage is estimated at runtime using an approximate algorithm to provide necessary memory information for subsequent task scheduling.

In data flow analysis, *memory manager* begins by traversing all kernel calls in the program, examining GPU-related pointer variables within the kernel parameters. Since kernels can only access memory in GPU space, these pointers should refer to specific GPU memory regions, each representing a distinct memory object. Notably, pointer aliasing can cause multiple pointers to reference the same memory region, so we trace allocation and deallocation instructions for each memory address by leveraging use-def chains. We treat any memory region allocated by the same allocation instruction as a memory object, and all kernels that use the pointers referencing

**Algorithm 2:** Instruction Scheduling Algorithm for Memory Lifetime Management

---

**Input:** Stream graph before transformation  $graph_{in}$   
**Output:** Stream graph after transformation  $graph_{out}$

```

1 Function PostponeMalloc( $graph$ ):
2    $memObjList \leftarrow graph.getMemObjList();$ 
3   for  $memObj \in memObjList$  do
4      $instrList \leftarrow graph.getRelatedInstr(memObj);$ 
5      $invokeList \leftarrow graph.getAssociatedKernels(memObj);$ 
6     sortByExecutionOrder( $invokeList$ );
7      $insertPoint \leftarrow invokeList[0];$ 
8      $graph.moveBefore(instrList, insertPoint);$ 
9     for  $i \leftarrow 1 \sim invokeList.size() - 1$  do
10       |  $graph.insertSyncBetween(insertPoint, invokeList[i]);$ 
11     end
12   end
13    $graph.removeRedundantSync();$ 
14   return  $graph;$ 
15 end

```

---

this region are considered dependent on that object. Data flow analysis enables *memory manager* to establish the dependency relationships between kernels and associated memory objects.

Once the dependencies between kernels and memory objects are determined, *memory manager* analyzes the memory objects' live ranges and schedules memory allocation and deallocation instructions to minimize their lifetimes. Algorithm 2 details the workflow of instruction scheduling to postpone GPU memory allocation. The algorithm takes the stream graph as input, which is generated by the stream scheduler and processed through data flow analysis. We assume that each memory object requires at most a single data transfer between host and device, as handling multiple transfers would necessitate a more sophisticated analysis to determine the optimal instruction placement. The algorithm begins by retrieving all memory objects within the program. For each memory object  $memObj$ , *memory manager* collects all allocation instructions and operations that may modify its content, such as data transfers or value assignments, and stores them in  $instrList$  in execution order for delayed allocation (line 4). Next, *memory manager* identifies the list of kernel calls that depend on  $memObj$ , known as  $invokeList$ , representing the object's live range (line 5). To determine the start of  $memObj$ 's live range,  $invokeList$  is sorted based on the original sequential execution order. The earliest kernel call is then identified as the beginning of the live range and serves as the insertion point for the associated instructions (line 6 ~ 7). *Memory manager* then moves  $instrList$  before the insertion point, and converts memory allocation and data transfer instructions to their asynchronous versions (e.g., `cudaMalloc` to `cudaMallocAsync` and `cudaMemcpy` to `cudaMemcpyAsync`) and assigns them to execute within the same stream as the insertion point (line 8). To ensure that  $memObj$  is allocated prior to any kernel call that depends on it, synchronization instructions are added between the insertion point and subsequent kernel calls (line 9 ~ 11). Finally, all redundant synchronization instructions within and across streams are removed, then returning the optimized stream graph (line 13 ~ 14). The algorithm for postponing free operations is similar to this algorithm and is not elaborated here.

After memory management, the peak memory usage of a task cannot be simply obtained by summing the sizes of all memory objects. To address this, we design an efficient algorithm for *lazy engine* to predict the peak memory usage during runtime. Given a stream graph with recorded operations, *lazy engine* first retrieves the delta memory




Fig. 8. Implementation of HUNTKTm. The compiler consists of several LLVM passes: *stream scheduler*, *memory manager*, *resource analyzer* and *function wrapper*, which transform the input code into a memory-optimized multi-stream program. During running, *lazy engine* maintains a CUDA operation queue and analyzes the task’s resource requirements. *Task dispatcher* manages a task queue and dispatches tasks to suitable devices. Each GPU has multiple concurrent streams for the parallel execution of kernels, along with dynamically sized memory pools for each task.

list from each stream, which logs memory changes of sequential memory allocation and deallocation. Then it accumulates these changes and takes the maximum value as the memory requirement for the stream. Summing these maximum values across all streams achieves peak memory usage of the stream graph. Note that our algorithm provides an upper bound of peak memory usage with  $O(N)$  time complexity, where  $N$  is the number of memory objects, since certain memory peaks may not occur due to inter-stream synchronization. Compared to our algorithm, accurate memory prediction would require enumerating numerous execution combinations and verifying inter-stream synchronization rules, which becomes computationally expensive as the number of memory objects increases.

### 5.5 Implementation

As shown in Figure 8, we implemented HUNTKTm on the basis of CUDA Runtime and LLVM Compiler Infrastructure [20]. Although targeting the CUDA platform, our design can be easily applied to other frameworks that support concurrent task queues and asynchronous memory management (e.g., HIP [2] and SYCL [11]). We pinpoint the pattern that kernels are always called by `__cudaPushCallConfiguration`, to find the serially issued kernels in host IR and apply our optimizations to their caller functions. To distinguish writable parameters from read-only ones, developers necessitate adding a parameter at the beginning of the kernel function’s parameter list, which indicates that the first  $N_{out}$  parameters are writable, and rearranging the writable parameters to the first  $N_{out}$  positions. This technique enables *DFG constructor* to analyze dependencies automatically, without involving any new programming framework.

To intercept CUDA runtime function calls and retrieve memory information, all memory-related function calls like `cudaMallocAsync` and `cudaFreeAsync` are wrapped by *function wrapper*. Similar transformations are applied to kernel launches for computational resource collection. During task execution, the program invokes wrapped functions to perform CUDA operations. *Lazy engine* analyzes and stores the call information in a queue, deferring execution until the task is dispatched to a specific device. When the program reaches `cudaTaskSchedule` or `cudaTaskScheduleLazy`, *lazy engine* sends the resource requirements to *task dispatcher* and waits for the device ID to be returned. The communication between *lazy engine* and *task dispatcher* is transferred over shared memory. A task is bound to the target device by calling `cudaSetDevice` after scheduling. For memory pool management in each task, HUNTKTm calls `cudaDeviceGetDefaultMemPool` to obtain the default memory pool before executing

the deferred operations and uses `cudaMemPoolSetAttribute` to set the memory release threshold to the predicted memory footprint. This attribute prevents memory from being released until usage exceeds the preset threshold.

## 6 Evaluation

### 6.1 Environment Setup

**6.1.1 Platform.** We conduct experiments on a server equipped with 4 NVIDIA A100 GPUs, 2 AMD EPYC 7742 64-Core Processors and 256 GB DDR4 memory. Each A100 GPU has 40 GB HBM and 6912 CUDA cores. The operating system is Debian 10.2.1 and the version of the NVIDIA driver is 555.42.06. We compile GPU programs using LLVM 14.0.6 and CUDA 12.4.0.

**6.1.2 Benchmark.** We use seven representative applications as listed in Table 2 to evaluate kernel-level parallelism. The two in-house micro-benchmarks are drawn from the kernels in NVIDIA FasterTransformer [31], and the rest of the benchmarks represent typical GPU workloads (image processing, machine learning, etc.), which are aligned with the benchmarks in GrSched [38]. Each application maintains multiple dependent kernels, some of which can be optimized to execute concurrently and overlap computation and data transfer to achieve higher performance.

We mark applications with memory requirements between 3 GB and 8 GB as small benchmarks (VEC, ML, DL) and requirements over 8 GB as large benchmarks (M1, M2, B&S, IMG). Benchmarks with various memory footprints are mixed in our workloads W1 to W8 with 4 different "large:small" ratios: 1:1, 2:1, 3:1, 5:1, similar to previous work [4]. We prefer using larger benchmarks to emulate the execution traces of heavy and long-running tasks in real-world workloads. Each workload consists of 16 or 32 tasks, randomly selected in proportion from small and large benchmark sets. All tasks within a workload arrive simultaneously and are scheduled as a single batch. The scheduler processes the batch by dequeuing one task and dispatching it to an appropriate device at a time, until the batch is empty or all devices are fully occupied.

Table 2. Evaluated benchmarks.

Name	Notation	DFG Width	Memory (GB)
Vector Square	VEC	2	4.80
Black & Scholes	B&S	10	12.8
Machine Learning	ML	2	3.12
Image Processing	IMG	3	11.6
Deep Learning	DL	2	7.06
Micro-1	M1	8	19.2
Micro-2	M2	6	17.6

Table 3. Workload mixes.

Workload	Jobs	Mix Ratio
W1	16	1:1
W2	16	2:1
W3	16	3:1
W4	16	5:1
W5	32	1:1
W6	32	2:1
W7	32	3:1
W8	32	5:1

**6.1.3 Evaluated Schemes.** In concurrent task execution scenarios, we compare HUNTKTm with two task scheduling designs: *single-assignment* (SA) [44], CASE [4]. SA assigns one job to each device at a time, and guarantees no device is idle when unhandled jobs exist. CASE automatically analyzes resource requirements of each task and schedules them according to available resources. For HUNTK, we extend SA by performing static stream scheduling, enabling multiple kernels within an application to execute concurrently in a device. HUNTKT incorporates stream scheduler and task scheduler to achieve hybrid scheduling, but lacks memory management. HUNTKTm is the complete design by integrating both hybrid scheduling and memory management to maximize kernel concurrency. NVIDIA MPS is enabled in both single and multiple GPU systems so that kernels from different processes can co-execute on the same device. To avoid exceeding the device's concurrent capacity,




Fig. 9. The throughput improvements of different task scheduling schemes across various workloads. The y-axis represents the throughput improvement in the multi-GPU system compared to the serial execution baseline SA.

HUNTKTM limits the number of available hardware queues per GPU to 32 in task scheduling algorithm, matching the maximum number of connections that the CUDA runtime can handle. Meanwhile, NVIDIA persistence mode [34] is enabled to reduce the GPU initialization overhead across applications.

For single task execution, we compare HUNTKTM against the baseline (serial execution, named *Serial* below), and two prior arts, including a static scheduler *Taskflow* [15] and a dynamic scheduler [38] based on GrCUDA [29] (denoted as *GrSched*). The maximum number of streams is set to 10 for both HUNTK and HUNTKTM, corresponding to the maximum DFG width among the benchmarks.

## 6.2 System Throughput

We first evaluate the system throughput of different scheduling schemes in task-concurrent scenarios across various workloads. As shown in Figure 9, HUNTKTM delivers performance improvements over other schemes in most workloads. This is mainly because HUNTKTM enables more kernels to run concurrently on the same device while ensuring load balancing across devices, effectively utilizing computational and memory resources. *CASE* exploits task-level concurrency to overlap computation and communication, yielding a 2.02x throughput gain over SA. HUNTK improves computational resource utilization by allowing multiple kernels within a single application to run concurrently. However, due to the limited opportunities for intra-application kernel concurrency, it achieves only a 1.20x average throughput improvement compared to serial execution. HUNTKT combines intra-task and inter-task kernel concurrency, while its task scheduler considers heterogeneous resource demands to enhance load balancing across devices, resulting in a 2.47x speedup. Based on HUNTKT, the complete design HUNTKTM further reduces the memory usage of applications, enabling more applications to run concurrently on the same device and efficiently utilizing idle resources. Ultimately, HUNTKTM achieves a 2.69x and 1.33x average performance improvement over SA and *CASE*, respectively.

As the proportion of large benchmarks increases and the number of tasks grows, memory becomes a bottleneck for concurrent execution, where *CASE* and HUNTKT are only able to run a limited number of benchmarks simultaneously. HUNTKTM addresses this issue by reducing memory requirements, allowing more tasks to run under the same memory capacity and achieving significant performance gains. For workload W1, which includes 16 applications with a 1:1 memory ratio, HUNTKTM and HUNTKT show similar speedups since the benefits of HUNTKTM are constrained by the number of benchmarks. In this case, all applications can be dispatched to devices by HUNTKTM without exceeding memory constraints. Meanwhile, as the number of concurrent tasks increases, memory management provides higher scheduling flexibility, leading HUNTKTM to deliver greater performance gains compared to *CASE*.

## 6.3 Hardware Resources Utilization

To analyze the impact of hybrid scheduling on system hardware resource utilization, we use DCGM [32], a low-overhead GPU system monitoring tool, to periodically collect hardware metrics. Two workloads W4 and W8




Fig. 10. Hardware metrics improvement achieved by different scheduling schemes and optimizations over W4 and W8 workloads.

Table 4. Memory consumption (GB) w/o and w/ memory management for mixed workloads.

Workload	HUNTKT	HUNTKTm	Reduction
W1	166.5	127.8	23.2%
W2	189.2	143.0	24.4%
W3	205.6	157.1	23.6%
W4	232.3	173.9	25.1%
W5	333.0	255.6	23.2%
W6	378.4	286.0	24.4%
W7	411.2	314.2	23.6%
W8	464.6	347.8	25.1%

Table 5. Memory consumption (GB) w/o and w/ memory management for applications.

Application	HUNTKT	HUNTKTm	Reduction
VEC	4.80	4.80	0%
B&S	12.8	12.8	0%
ML	3.12	3.08	1.3%
IMG	11.6	9.20	20.0%
DL	7.06	4.70	33.3%
M1	19.2	13.4	30.0%
M2	17.6	11.2	36.4%
Avg.	10.9	8.47	22.3%

are selected for analysis as they demonstrate the task scheduling efficiency under GPU memory constraints, and the results are shown in Figure 10. The results show that HUNTKT is more effective than HUNTK in utilizing idle resources. This is because the performance of HUNTK is limited by the number of kernels that can execute concurrently within an application and the proportion of kernel execution time relative to overall task time. By parallelizing computation and data communication across multiple independent tasks, HUNTKT achieves higher resource utilization than HUNTK. Leveraging kernel-level concurrency within applications in a task-concurrent environment, HUNTKT improves FP32 utilization, memory bandwidth utilization, and SM occupancy by 3.54x, 2.83x, and 2.47x on average under W4 and W8 compared to SA, significantly outperforming HUNTK and CASE, which rely solely on intra- or inter-application concurrent kernel execution. With memory management enabled, HUNTKTm achieves even greater improvements of 4.45x, 3.39x, and 3.76x, corresponding to utilization gains of 91.0%, 45.5%, and 111.2% over CASE, demonstrating its ability to further enhance resource utilization through memory optimization in highly concurrent environments. For workloads W4 and W8, HUNTKTm significantly increases SM occupancy and FP32 utilization, leveraging more idle computing resources to improve the computational efficiency of the system. The improvement in bandwidth utilization highlights the ability of HUNTKTm to overlap computation and communication when allowing more tasks to execute simultaneously, alleviating the inefficiencies caused by serial execution of computation and data transfers within applications.

## 6.4 Memory Reduction

**6.4.1 Mixed Workload.** Table 4 summarizes the cumulated memory requirements for each workload before and after applying memory management in HUNTKTm. By leveraging memory reuse based on liveness analysis, HUNTKTm achieves a significant reduction in memory consumption, lowering the total memory usage by 23.2% to 25.1%. With reduced memory requirements, HUNTKTm enables the simultaneous execution of all benchmarks




Fig. 11. Average speedup gained by different schemes in seven applications.

in workloads W1 and W2, where the number of tasks in the workload becomes the limiting factor for further kernel concurrency. While the memory footprint of W3 is reduced to 157.1 GB after memory optimization, which is slightly below the system GPU memory capacity, unavoidable memory fragmentation prevents the immediate execution of the final task in W3. In workload W8, which contains the largest number of benchmarks, HUNTKTm is able to launch 14 tasks simultaneously, whereas HUNTKT supports the concurrent execution of only 9 tasks. Concurrent execution of additional tasks maximizes hardware resource utilization while effectively overlapping computation and data transfer operations. Therefore, the kernel concurrency improvement achieved through memory management enhances system throughput, with these benefits being particularly pronounced in scenarios with high memory demands.

**6.4.2 Single Benchmark.** Table 5 compares memory consumption before and after applying memory management in HUNTKTm. For most evaluated benchmarks, HUNTKTm effectively reduces the peak memory usage by an average of 22.3%, which is calculated as a weighted average of individual reduction ratios, with each benchmark weighted by its original peak memory usage. The ratio of memory reduction achieved by HUNTKTm depends on the data dependencies between kernels within the application. We define the kernel execution path as the number of kernels executed sequentially within a single stream. Generally, longer execution paths provide more opportunities for HUNTKTm to optimize memory usage by shortening the lifetime of memory objects, as more allocation and deallocation operations can be scheduled within a single stream. The execution path of M2, which includes several accumulation and activation kernels, achieves a memory reduction of 36.4% by releasing unused memory after each kernel. However, HUNTKTm does not always lead to memory savings. ML includes multiple in-place operators whose inputs cannot be released after kernel execution, limiting HUNTKTm to reusing smaller memory objects. Additionally, to maximize computational efficiency, HUNTKTm distributes independent kernels across multiple streams, which may shorten the execution paths. For example, VEC’s three kernels are executed across two streams, and B&S’s ten kernels are evenly distributed across ten streams, leaving no opportunities for memory reuse.

## 6.5 Task-level Performance Improvement

In this section, we evaluate the GPU operation execution time for the selected programs. Figure 11 presents the speedup achieved by various kernel-level concurrency schemes. *GrSched*, which leverages unified memory, suffers from substantial overhead during data transfers, leading to execution times an order of magnitude slower than *Serial*. Across the evaluated benchmarks, *Taskflow*, HUNTKT, and HUNTKTm achieve average speedups of 1.67x, 1.69x, and 1.90x, respectively. While both HUNTKT and *Taskflow* allocate kernels to different streams, CUDA Graph construction and initialization overhead slightly hinder *Taskflow*’s performance compared to HUNTKT, particularly for applications with short kernel execution times. HUNTKTm improves performance by introducing memory pools to minimize the overhead of frequent memory allocations and deallocations. It also schedules memory operations closer to kernel launches, enabling better overlap of computation and communication. As a

Table 6. Code modification based on serial version (LoC / # tokens).

Scheme	VEC	B&S	ML	IMG	DL	M1	M2	Avg.
Async	19/137	40/349	30/209	45/394	29/267	67/588	51/443	40/341
Taskflow	11/136	18/363	28/414	26/445	33/421	39/651	34/569	27/428
GrSched	56/415	116/1109	109/1061	153/1515	146/1404	159/1716	147/1582	127/1257
CASE	0/0	0/0	0/0	0/0	0/0	0/0	0/0	0/0
HUNTKTm	5/17	11/34	19/66	19/65	13/44	20/63	17/54	15/49

result, HUNTKTm improves speedups to 3.27x and 3.17x for M1 and M2, respectively. However, in some cases (e.g., B&S), the computation time within each stream is shorter than the data transfer time. The data transfer operations distributed around kernel launches revert concurrent kernels to serial execution, yielding a modest performance improvement of 1.10x.

We also compare the kernel execution time achieved by different schemes. Since HUNTKTm interleaves memory-related instructions between kernels, it is excluded from the comparison. Similar to Figure 11, HUNTKT and *Taskflow* achieve comparable speedups in most benchmarks, averaging 2.79x and 2.99x. For *GrSched*, the most significant performance penalty comes from dynamic scheduling, including runtime dependency analysis, kernel capture, and kernel issue. Furthermore, this overhead prevents *GrSched* from launching multiple kernels concurrently, limiting its ability to achieve optimal kernel-level concurrency. As a result, *GrSched* achieves an average speedup of 1.92x, significantly lower than HUNTKT, which leverages static analysis and optimization.

## 6.6 Programming Efforts

Table 6 lists the programming efforts required by different concurrent schemes based on serial version programs, in terms of the average modified lines of code (LoC) and token count. *Async* means programming with expert concurrency optimization using CUDA’s APIs. Compared with *Serial*, HUNTKTm costs only 15 LoC and 49 tokens in addition to enabling CKE and automating both dependency analysis and memory management. The extra code is sourced from the kernels’ light-weight wrapper for writable parameter identification, and our compiler-based approach encompasses the rest of the transformation and optimization. In contrast, significant code modification is involved in other schemes. *Async* necessitates manually managing CUDA’s asynchronous APIs, including stream initialization, synchronization and distributing kernels to streams, which require tremendous programming efforts and are error-prone. *Taskflow* lessened this burden and still requires explicit dependence specification. *GrSched* has the merit of automation but is limited to a dynamic programming language for runtime analysis. By automating dependency analysis and memory management, HUNTKTm eliminates common errors such as incorrect stream synchronization and misordered memory operations, which are difficult to debug manually.

Unlike other schemes, *CASE* does not require source code modification as it only retrieves the program’s resource requirements without considering dependencies between kernels and relies on the runtime system for task scheduling. In comparison to CASE, HUNTKTm introduces lightweight code modifications, primarily a one-line parameter addition for each kernel definition and launch, while providing support for both dependency-aware kernel scheduling and memory optimization. This trade-off is especially valuable for large-scale or performance-critical applications where minor code adjustments are acceptable for substantial runtime gains.

## 6.7 Overhead

**6.7.1 Compilation.** The static compilation overhead comprises three components: kernel scheduling, resource analysis, and memory management. All costs arise from code analysis and transformation without kernel profiling. The execution time of kernel scheduling and memory management depends on the number of kernel invocations and memory objects, respectively. Resource analysis overhead includes retrieving register and shared memory




Fig. 12. Average throughput improvement of different task scheduling schemes when limiting GPUs' memory capacities to 20GB, 30GB and 40GB.




Fig. 13. Speedup and memory ratio achieved by HUNTKTm when the number of streams ranges from 1 to 10.

usage from nvcc [27] and analyzing the arguments of kernel launches and memory allocations. Across seven benchmark applications, the average compilation time rose from 1.41 s to 2.40 s, which is an acceptable cost in compilation.

**6.7.2 Runtime.** At runtime, the overhead mainly consists of task scheduling and kernel launch preparation. Task scheduling involves traversing the stream graph to estimate peak memory usage and iterating over available devices, which takes around 1 millisecond and is thus negligible. Before launching kernels, several preparation steps are invoked to ensure concurrent execution, including CUDA APIs such as stream creation, event synchronization and the CUDA context initialization. The primary runtime overhead stems from CUDA context initialization, whose average time increases from 81 ms to 118 ms due to multiple tasks sharing a single MPS server. However, context initialization can be overlapped with the computation and communication of other tasks, incurring no additional execution time.

## 6.8 Sensitivity Studies

**6.8.1 Memory Capacity.** In task-concurrent scenarios, memory capacity is a critical factor limiting the number of tasks that can run simultaneously on a device. To evaluate the schedulers' performance under different memory constraints, we constrain the memory capacity of each GPU in the multi-GPU system to 20GB, 30GB and 40GB. Figure 12 presents the system throughput under various scheduling schemes, which are normalized to SA. As memory capacity decreases, many workloads fail to launch due to insufficient memory, causing a significant decline in throughput for CASE. HUNTKT mitigates idle computational resources by issuing more kernels from a single application across multiple hardware queues. However, the performance gains of hybrid scheduling remain constrained by memory bottlenecks. HUNTKTm reduces the memory usage of most tasks by memory reusing, enabling more tasks to collocate in the same device with limited memory. As a result, HUNTKTm achieves throughput improvements of 2.12x and 2.38x under 20 GB and 30 GB memory constraints, respectively, with only a slight decrease compared to the 2.69x speedup achieved with 40 GB memory. Compared to CASE, HUNTKTm outperforms by 61.8%, 51.6% and 33.2% on average under three memory capacities, demonstrating that hybrid scheduling and memory management deliver significant performance gains under varying memory constraints.

**6.8.2 Number of Streams.** In the design of HUNTKTm, the number of streams plays a critical role in determining program concurrency performance and memory usage. Figure 13(a) shows the execution speedup of three programs under HUNTKTm compared to serial execution as the number of streams increases from 1 to 10. We select B&S, M1, and M2 as the subjects of the experiment as they exhibit substantial kernel-level concurrency.




Fig. 14. Average throughput improvement of different task scheduling schemes when the number of GPUs ranges from 1 to 4 on A100 or 4090 systems.

When the number of streams is 1, three applications achieve an average speedup of 1.12x due to the memory pool mechanism, which eliminates frequent allocation and deallocation operations. M1 and M2 achieve their peak speedups of 3.44x and 3.18x when the number of streams matches their DFG widths, enabling maximal kernel-level parallelism within each task. Beyond these points, increasing the number of streams yields no further performance gains. Since data transfers without pinned memory cannot execute concurrently, frequent serialized transfers in B&S block nearly all operations within the application from executing concurrently, even when a sufficient number of streams are available.

The impact of the number of streams on memory management is shown in Figure 13(b). As the number of streams decreases, memory usage for the three benchmarks is reduced, with maximum memory savings of 90.0%, 63.3%, and 59.0% achieved under single stream execution. Similar to the acceleration from multi-stream concurrency, when the number of streams equals the DFG width, all independent kernels are assigned to separate streams. At this point, memory reuse within streams depends entirely on the execution paths of dependent kernels, resulting in memory reductions of 0%, 30%, and 36%, respectively. We observe that each stream in B&S contains only one kernel, and thus no memory reuse opportunities exist when using 10 streams. Therefore, users should choose an appropriate number of streams to balance execution efficiency and memory usage for programs.

**6.8.3 Number of GPUs.** Figure 14(a) illustrates the average throughput achieved by different task scheduling methods as the number of GPUs increases from 1 to 4 on the A100 system. When the number of available GPUs is limited, a few benchmarks with large memory footprints dominate the devices, restricting throughput gains by reducing task concurrency. CASE and HUNTK employ naive kernel-level or task-level concurrency, allowing only a limited number of kernels to execute simultaneously, which fails to fully utilize GPU hardware resources. In contrast, HUNTKTm leverages hybrid scheduling and memory management to significantly increase kernel concurrency, achieving throughput improvements of 4.42x, 3.68x and 2.88x over SA when the number of GPUs ranges from 1 to 3, respectively. As GPUs increase, extensive data transfers between the host and devices emerge as a bottleneck, limiting further performance gains. The increased memory transfer time forces some potentially parallel kernels to execute sequentially, reducing the concurrency benefits of HUNTKTm. In our future work, we will focus on incorporating memory bandwidth requirements of concurrent tasks into the task scheduler to address this limitation.

**6.8.4 Hardware Platform.** To evaluate the performance of HUNTKTm across different hardware platforms, we conduct experiments on a system equipped with 4 NVIDIA RTX 4090 24GB GPUs, 2 Intel Xeon Gold 6338N CPUs, and 1024 GB DRAM. Figure 14(b) shows the throughput improvements of various workloads on the 4090 system. Similar to A100 system, the throughput improvement of HUNTKTm over SA increases as the number of GPUs decreases, reaching 4.12x, 3.93x, 3.60x, and 3.42x for 1 to 4 GPUs, respectively. Since RTX 4090 has lower memory bandwidth compared to A100, longer data transfer times are overlapped with computation across tasks. In addition, the higher number of SMs and CUDA cores in the 4090 enables greater kernel-level concurrency. These




Fig. 15. Average speedup gained by different schemes in seven applications when kernel execution times become extremely short.

factors contribute to the more pronounced performance gains of HUNTKTM on the 4090 system. As previously discussed, HUNTKTM demonstrates greater advantages under memory-constrained conditions. On the 4090 system, which provides only 24GB memory per GPU, HUNTKTM achieves an average performance improvement of 52.5% over CASE across different GPU numbers.

**6.8.5 Short Kernel Execution Times.** To further analyze the impact of short kernel execution times on task-level performance under different kernel concurrency strategies, we reduce the input size of each application to one-thousandth of its original scale, and the results are shown in Figure 15. When kernel execution time is extremely short, HUNTKT still achieves an average 1.35x speedup over the *Serial*, as it distributes kernels across streams without introducing additional overhead. In contrast, HUNTKTM incurs overhead due to its memory pool mechanism, where asynchronous memory allocation and deallocation introduce additional costs. For example, in VEC, the memory allocation time increases from 250 us to 860 us. While this overhead is negligible in typical scenarios, it becomes significant with short-running kernels, leading to performance degradation in VEC, B&S, and ML. As a result, HUNTKTM achieves an average 1.21x speedup across all benchmarks. On the other hand, *Taskflow* introduces substantial runtime overhead from thread creation, synchronization, and destruction, along with non-trivial CUDA graph creation costs, ultimately causing a severe deterioration in performance and yielding only a 0.54x speedup.

## 7 Conclusion

This paper introduces HUNTKTM, a hybrid scheduling and automatic management framework designed to optimize both kernel-level and task-level concurrency for efficient GPU execution. HUNTKTM integrates three key components - a stream scheduler, a task scheduler, and a memory manager - to form a unified execution stack. The stream scheduler identifies kernel dependencies and distributes kernels across multiple concurrent streams. The task scheduler analyzes resource requirements and dynamically dispatches tasks to appropriate devices. Evaluations show that HUNTKTM achieves substantial performance gains in both task-concurrent and multi-kernel execution scenarios, delivering an average 33.2% speedup over CASE and 13.8% over Taskflow, respectively. These improvements highlight the effectiveness of HUNTKTM’s coordinated approach to scheduling and memory management. In the future, we plan to extend HUNTKTM to support highly distributed, multi-node GPU systems, and further refine its scheduling and management strategies to scale across a broader range of complex applications.

## References

- [1] Martin Abadi, Paul Barham, Jianmin Chen, Zhifeng Chen, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Geoffrey Irving, Michael Isard, et al. 2016. TensorFlow: a system for Large-Scale machine learning. In *12th USENIX symposium on operating systems design and implementation*. 265–283.
- [2] AMD. 2016. HIP:Heterogeneous Interface for Portability. <https://github.com/ROCm-Developer-Tools/HIP>.

- [3] Mehmet E Belviranli, Seyong Lee, Jeffrey S Vetter, and Laxmi N Bhuyan. 2018. Juggler: a dependence-aware task-based execution framework for GPUs. In *Proceedings of the 23rd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*. 54–67.
- [4] Chao Chen, Chris Porter, and Santosh Pande. 2022. Case: A compiler-assisted scheduling framework for multi-gpu systems. In *Proceedings of the 27th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*.
- [5] Guoyang Chen, Yue Zhao, Xipeng Shen, and Huiyang Zhou. 2017. Effisha: A software framework for enabling efficient preemptive scheduling of gpu. In *Proceedings of the 22nd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*. 3–16.
- [6] Quan Chen, Hailong Yang, Minyi Guo, Ram Srivatsa Kannan, Jason Mars, and Lingjia Tang. 2017. Prophet: Precise qos prediction on non-preemptive accelerators to improve utilization in warehouse-scale computers. In *Proceedings of the Twenty-Second International Conference on Architectural Support for Programming Languages and Operating Systems*. 17–32.
- [7] Tianqi Chen, Bing Xu, Chiyuan Zhang, and Carlos Guestrin. 2016. Training deep nets with sublinear memory cost. *arXiv preprint arXiv:1604.06174* (2016).
- [8] Weiduo Chen, Xiaoshe Dong, Fan Zhang, Bowen Li, Yufei Wang, and Qiang Wang. 2024. ATP: Achieving Throughput Peak for DNN Training via Smart GPU Memory Management. *ACM Transactions on Architecture and Code Optimization* (2024).
- [9] Seungbeom Choi, Sunho Lee, Yeonjae Kim, Jongse Park, Youngjin Kwon, and Jaehyuk Huh. 2022. Serving heterogeneous machine learning models on Multi-GPU servers with Spatio-Temporal sharing. In *2022 USENIX Annual Technical Conference*. 199–216.
- [10] Aditya Dhakal, Sameer G Kulkarni, and KK Ramakrishnan. 2020. Gslice: controlled spatial sharing of gpus for a scalable inference platform. In *Proceedings of the 11th ACM Symposium on Cloud Computing*. 492–506.
- [11] The Khronos SYCL Working Group. 2020. SYCL 2020 Specification. <https://registry.khronos.org/SYCL/specs/sycl-2020/pdf/sycl-2020.pdf>.
- [12] Vishakha Gupta, Karsten Schwan, and Niraj Tolia et al. 2011. Pegasus: Coordinated Scheduling for Virtualized Accelerator-based Systems. In *USENIX Annual Technical Conference*. Portland, OR, USA.
- [13] Mingcong Han, Hanze Zhang, Rong Chen, and Haibo Chen. 2022. Microsecond-scale preemption for concurrent GPU-accelerated DNN inferences. In *16th USENIX Symposium on Operating Systems Design and Implementation*. 539–558.
- [14] Qingda Hu, Jiwu Shu, Jie Fan, and Youyou Lu. 2016. Run-time performance estimation and fairness-oriented scheduling policy for concurrent GPGPU applications. In *45th International Conference on Parallel Processing*. IEEE.
- [15] Tsung-Wei Huang, Dian-Lun Lin, Chun-Xun Lin, and Yibo Lin. 2021. Taskflow: A lightweight parallel and heterogeneous task graph computing system. *IEEE Transactions on Parallel and Distributed Systems* (2021), 1303–1320.
- [16] Yafan Huang, Sheng Di, Guanpeng Li, and Franck Cappello. 2024. cuSZp2: A GPU Lossy Compressor with Extreme Throughput and Optimized Compression Ratio. In *2024 SC24: International Conference for High Performance Computing, Networking, Storage and Analysis SC*. IEEE Computer Society, 188–205.
- [17] Paras Jain, Ajay Jain, Aniruddha Nrusimha, Amir Gholami, Pieter Abbeel, Joseph Gonzalez, Kurt Keutzer, and Ion Stoica. 2020. Checkmate: Breaking the memory wall with optimal tensor rematerialization. *Proceedings of Machine Learning and Systems* (2020), 497–511.
- [18] Jaehoon Jung, Jinpyo Kim, and Jaejin Lee. 2023. DeepUM: Tensor migration and prefetching in unified memory. In *Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*. 207–221.
- [19] Yunseong Kim, Yujeong Choi, and Minsoo Rhu. 2022. Paris and elsa: An elastic scheduling algorithm for reconfigurable multi-gpu inference servers. In *Proceedings of the 59th ACM/IEEE Design Automation Conference*. 607–612.
- [20] C. Lattner and V. Adve. 2004. LLVM: a compilation framework for lifelong program analysis & transformation. In *International Symposium on Code Generation and Optimization*. 75–86.
- [21] Jaeho Lee, Shinnung Jeong, Seungbin Song, Kunwoo Kim, Heelim Choi, Youngsok Kim, and Hanjun Kim. 2023. Occamy: Memory-efficient GPU Compiler for DNN Inference. In *60th ACM/IEEE Design Automation Conference*. IEEE.
- [22] Ruoxiang Li, Tao Hu, Xu Jiang, Laiwen Li, Wenxuan Xing, Qingxu Deng, and Nan Guan. 2023. Rosgm: A real-time gpu management framework with plug-in policies for ros 2. In *2023 IEEE 29th Real-Time and Embedded Technology and Applications Symposium*. IEEE, 93–105.
- [23] Xinjian Long, Xiangyang Gong, Bo Zhang, and Huiyang Zhou. 2023. Deep learning based data prefetching in CPU-GPU unified virtual memory. *J. Parallel and Distrib. Comput.* (2023).
- [24] Xinjian Long, Xiangyang Gong, Bo Zhang, and Huiyang Zhou. 2023. An intelligent framework for oversubscription management in cpu-gpu unified memory. *Journal of Grid Computing* (2023), 11.
- [25] Bernabé López-Albelda, Francisco M Castro, José M González-Linares, and Nicolás Guil. 2022. FlexSched: Efficient scheduling techniques for concurrent kernel execution on GPUs. *The Journal of Supercomputing* (2022), 43–71.
- [26] Abdun Nihaal and Madhu Mutyam. 2024. Selective Memory Compression for GPU Memory Oversubscription Management. In *Proceedings of the 53rd International Conference on Parallel Processing*. 189–198.
- [27] NVIDIA. 2007. CUDA LLVM Compiler. <https://developer.nvidia.com/cuda-llvm-compiler>.
- [28] NVIDIA. 2017. NVIDIA Multi-Process Service (MPS). <https://docs.nvidia.com/deploy/mps/index.html>.
- [29] NVIDIA. 2020. grCUDA: Polyglot GPU Access in GraalVM. <https://github.com/NVIDIA/grcuda>.
- [30] NVIDIA. 2020. NVIDIA Multi-Instance GPU (MIG). <https://www.nvidia.com/en-us/technologies/multi-instance-gpu>.
- [31] NVIDIA. 2021. Faster transformer. <https://github.com/NVIDIA/FasterTransformer>.

- [32] NVIDIA. 2021. Manage and Monitor GPUs in Cluster Environments. <https://developer.nvidia.com/dcgm>.
- [33] NVIDIA. 2023. CUDA C++ Programming Guide. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>.
- [34] NVIDIA. 2024. NVIDIA Driver Persistence. <https://docs.nvidia.com/deploy/driver-persistence/index.html>.
- [35] OpenMP. 2023. OpenMP. <https://www.openmp.org/>.
- [36] Sreepathi Pai, Matthew J. Thazhuthaveetil, and R. Govindarajan. 2013. Improving GPGPU Concurrency with Elastic Kernels. In *Proceedings of the Eighteenth International Conference on Architectural Support for Programming Languages and Operating Systems*. 407–418.
- [37] Jason Jong Kyu Park, Yongjun Park, and Scott Mahlke. 2015. Chimera: Collaborative preemption for multitasking on a shared GPU. *ACM SIGARCH Computer Architecture News* (2015), 593–606.
- [38] Alberto Parravicini, Arnaud Delamare, and Marco Arnaboldi et al. 2021. DAG-based Scheduling with Resource Sharing for Multi-task Applications in a Polyglot GPU Runtime. In *35th IEEE International Parallel and Distributed Processing Symposium*. Portland, OR, USA, 111–120.
- [39] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. 2019. Pytorch: An imperative style, high-performance deep learning library. *Advances in neural information processing systems* (2019).
- [40] Manos Pavlidakis, Giorgos Vasiliadis, Stelios Mavridis, Anargyros Argyros, Antony Chazapis, and Angelos Bilas. 2024. Guardian: Safe GPU Sharing in Multi-Tenant Environments. In *Proceedings of the 25th International Middleware Conference*. 313–326.
- [41] Sooraj Puthoor, Xulong Tang, Joseph Gross, and Bradford M Beckmann. 2018. Oversubscribed command queues in GPUs. In *Proceedings of the 11th Workshop on General Purpose GPUs*. 50–60.
- [42] Jiaxing Qi, Wencong Xiao, Mingzhen Li, Chaojie Yang, Yong Li, Wei Lin, Hailong Yang, Zhongzhi Luan, and Depei Qian. 2024. ElasticBatch: A Learning-Augmented Elastic Scheduling System for Batch Inference on MIG. *IEEE Transactions on Parallel and Distributed Systems* (2024).
- [43] Zhengwei Qi, Jianguo Yao, Chao Zhang, Miao Yu, Zhizhou Yang, and Haibing Guan. 2014. VGRIS: Virtualized GPU resource isolation and scheduling in cloud gaming. *ACM Transactions on Architecture and Code Optimization* (2014), 1–25.
- [44] Carlos Reano, Federico Silla, Dimitrios S Nikolopoulos, and Blession Varghese. 2017. Intra-node memory safe gpu co-scheduling. *IEEE Transactions on Parallel and Distributed Systems* (2017), 1089–1102.
- [45] Milan Shah, Xiaodong Yu, Sheng Di, Michela Becchi, and Franck Cappello. 2023. Lightweight Huffman Coding for Efficient GPU Compression. In *Proceedings of the 37th International Conference on Supercomputing*. 99–110.
- [46] S-Kazem Shekofteh, Hamid Noori, Mahmoud Naghibzadeh, Holger Fröning, and Hadi Sadoghi Yazdi. 2019. cCUDA: Effective co-scheduling of concurrent kernels on GPUs. *Transactions on Parallel and Distributed Systems* (2019), 766–778.
- [47] S-Kazem Shekofteh, Hamid Noori, Mahmoud Naghibzadeh, Hadi Sadoghi Yazdi, and Holger Fröning. 2019. Metric selection for GPU kernel classification. *ACM Transactions on Architecture and Code Optimization* (2019), 1–27.
- [48] Alex Suhan, Davide Libenzi, Ailing Zhang, Parker Schuh, Brennan Saeta, Jie Young Sohn, and Denys Shabalin. 2021. LazyTensor: combining eager execution with domain-specific compilers. *arXiv preprint arXiv:2102.13267* (2021).
- [49] Zhenbo Sun, Huanqi Cao, Yuanwei Wang, Guanyu Feng, Shengqi Chen, Haojie Wang, and Wenguang Chen. 2024. AdaPipe: Optimizing Pipeline Parallelism with Adaptive Recomputation and Partitioning. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*.
- [50] Ivan Tanasic, Isaac Gelado, Javier Cabezas, Alex Ramirez, Nacho Navarro, and Mateo Valero. 2014. Enabling preemptive multiprogramming on GPUs. *ACM SIGARCH Computer Architecture News* (2014), 193–204.
- [51] Yu Tang, Qiao Li, Lujia Yin, Dongsheng Li, Yiming Zhang, Chenyu Wang, Xingcheng Zhang, Linbo Qiao, Zhaoning Zhang, and Kai Lu. 2024. DELTA: Memory-Efficient Training via Dynamic Fine-Grained Recomputation and Swapping. *ACM Transactions on Architecture and Code Optimization* (2024).
- [52] Nandita Vijaykumar, Kevin Hsieh, Gennady Pekhimenko, Samira Khan, Ashish Shrestha, Saugata Ghose, Adwait Jog, Phillip B Gibbons, and Onur Mutlu. 2016. Zorua: A holistic approach to resource virtualization in GPUs. In *49th Annual IEEE/ACM International Symposium on Microarchitecture*. IEEE, 1–14.
- [53] Zhenning Wang, Jun Yang, Rami Melhem, Bruce Childers, Youtao Zhang, and Minyi Guo. 2016. Simultaneous multikernel GPU: Multi-tasking throughput processors via fine-grained sharing. In *2016 IEEE international symposium on high performance computer architecture*. IEEE, 358–369.
- [54] Yue Weng, Tianao Ge, Xi Zhang, Xianwei Zhang, and Yutong Lu. 2022. Raise: Efficient gpu resource management via hybrid scheduling. In *22nd IEEE International Symposium on Cluster, Cloud and Internet Computing*. IEEE, 685–695.
- [55] Bo Wu, Xu Liu, Xiaobo Zhou, and Changjun Jiang. 2017. Flep: Enabling flexible and efficient preemption on gpus. *ACM SIGPLAN Notices* (2017), 483–496.
- [56] Hao Wu, Weizhi Liu, Huanxin Lin, and Cho-Li Wang. 2020. A model-based software solution for simultaneous multiple kernels on GPUs. *ACM Transactions on Architecture and Code Optimization* (2020), 1–26.

- [57] Qiumin Xu, Hyeran Jeon, Keunsoo Kim, Won Woo Ro, and Murali Annavaram. 2016. Warped-slicer: Efficient intra-SM slicing through dynamic resource partitioning for GPU multiprogramming. *ACM SIGARCH Computer Architecture News* (2016), 230–242.
- [58] Su-Wei Yang, Zhao-Wei Qiu, and Ya-Shu Chen. 2020. GPU swap-aware scheduler: virtual memory management for GPU applications. In *Proceedings of the 35th Annual ACM Symposium on Applied Computing*. 1222–1227.
- [59] Tsung Tai Yeh, Amit Sabne, and Putt Sakdhnagool et al. 2017. Pagoda: Fine-Grained GPU Resource Virtualization for Narrow Tasks. In *Proceedings of the 22nd ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*.

Received 22 January 2025; revised 5 August 2025; accepted 16 October 2025

Just Accepted